

Advanced Programming Tutorial

Worksheet 4: Resource Management

Keywords: raw pointers, `new`, `delete`, address sanitizer, `unique_ptr`, `std::map`, `struct`, `->`

Exercise 1: Matrix-vector product - A Space Odyssey

Sometimes you can't escape your legacy. There will be times that you have to deal with lower-level code. Let's explore how we could have implemented the same matrix-vector product from Worksheet 3 in 2001.

In older C++, we used to represent 2D matrices as pointers to 1D arrays, like this:

```
1  double * matrix = new double[matrix_rows * matrix_columns];  
2  matrix[i*matrix_columns+j]
```

The file `matrix-operations-part3.cpp` contains the solution from Worksheet 3 and an example implementation `mat_times_vec()` of the matrix-vector product we implemented in `operator*`.

1. Read through the code: what differences do you observe? Which additional keywords, types, and variables do you see?
2. Important to remember is that we always need to `delete` everything that we allocate with `new`. However, it is easy to forget something. Compile your code with GCC and use the option `-fsanitize=address`. What additional information do you get at runtime?
3. There is a memory leak in the program: fix it!
4. **Idea:** Can you think of a better way to organize your code so that you can more easily see the potential memory leaks? Think of a principle we discussed in the lecture.

Note: You are of course using Git to keep track of your exercises, right? ;-) Since we will not go too far in this direction, you may want to make a separate git branch for this.

Exercise 2: The pandemic (tracker) gets real (data)

In Worksheet 2, you already created your first pandemic tracker. For the purpose of quick development, you used pseudo-randomly generated data. Now you want to also process real data from a CSV file¹ However, you want to keep the dummy data to be able to easily test your implementation, without changing the code.

In this exercise, we will use a `unique_ptr` to select between the dummy and the real dataset at runtime. We will also learn how to pass the `unique_ptr` to a function and we will work with `std::map` to create views of the dataset.

Open the skeleton file `vis_app_part2.cpp` and:

1. Take 5-10min to read through the code and the section after this exercise. There are code parts you don't need to read right now (see guiding comments).
2. Fill the TODO parts in `main()`:
 - (a) Make the unique pointer `data_frame` so that it points to dummy or real data, depending on the `if` branch.
 - (b) Print the cases for each country, independent of the kind of data (dummy or real).
3. Implement `normalize_per_capita(std::unique_ptr<DataFrame>& data_frame)`. This should return a map named `cases_normalized` with the cases per 100,000 people for each country and day.
4. Call the function `normalize_per_capita` on the given `data_frame` and print the results.

Some new concepts

Structures: struct

You may already know classes: that's good! You can think of structures as simplified classes (with public members by default). You don't already know what a class is? Then you can think of a struct as a collection of data, such as:

```
1 struct Color {  
2     int red;  
3     int green;  
4     int blue;  
5 };
```

We can then create a specific color and access the underlying elements as:

```
1 Color TUMBlue = {0, 101, 189};  
2 std::cout << TUMBlue.red << std::endl;
```

¹European Centre for Disease Prevention and Control, file formatted for our needs.

Arrow operator: ->

If we access our color through a pointer (raw or smart), we first need to dereference it to access the members of it:

```
1 Color TUMBlue = {0, 101, 189}; // ok, normally no temporary object needed
2 unique_ptr<Color> myFavoriteColor = make_unique<Color>(TUMBlue);
3 std::cout << (*myFavoriteColor).red << std::endl;
```

To make our lives easier, we can alternatively write the equivalent:

```
1 std::cout << myFavoriteColor->red << std::endl;
```

Maps: std::map

A map is a data structure that maps *keys* to *values*. For example:

```
1 Color TUMBlue = {0, 101, 189};
2 Color TUMGreen = {162, 173, 0};
3 Color TUMLightBlue = {100, 160, 200};
4 std::map<std::string, Color> TUMColorPalette {
5     {"color-logo", TUMBlue},
6     {"color-highlight", TUMGreen},
7     {"color-notes", TUMLightBlue}
8 }
```

To iterate over a map, we can use a ranged-for loop:

```
1 for( const auto & [name, color] : TUMColorPalette ) {
2     std::cout << "Name: " << name << ", So much red: " <<
3         << color.red << std::endl;
4 }
```

We can also directly access a map element like this:

```
1 TUMColorPalette["color-logo"] = TUMLightBlue;
```

Important: If "color-logo" does not already exist in the map, then attempting to access `TUMColorPalette["color-logo"]` (without assigning anything to it) would add an empty value to this key. To check if key exists, use `map.contains(key)` (C++20).

Advice / C++ Core Guidelines

- R.1: Manage resources automatically using resource handles and RAII
- R.3: A raw pointer (a T*) is non-owning
- R.5: Prefer scoped objects, don't heap-allocate unnecessarily
- R.10: Avoid `malloc()` and `free()`
- R.11: Avoid calling `new` and `delete` explicitly